

SKETCH · ARTWORK

Sketch in. Artwork out.

A web studio that redraws your sketch into finished art — a **Stable Diffusion img2img** engine behind an interface I designed and built solo.

UI/UX Design · HTML · CSS · JS · Stable Diffusion



THE BRIEF

It *develops your sketch*, not replaces it.

SketchMuse develops the sketch you started with, tuned over a thousand test renders.

ROLE
UI/UX · Frontend · Model integration

YEAR
2025 · Solo project

PLATFORM
Responsive web app

BUILT WITH
Figma · HTML/CSS/JavaScript
Stable Diffusion Model

PROJECT OVERVIEW

- 01 **The engine wasn't the problem**

Upload a sketch, and what came back was never mine. Only running the model myself could develop it further.
- 02 **What I built**


tuned the model, designed the interface and built the front end, taking the whole product from model to browser.
- 03 **What I found**

Five people tested it. The names scored 4.2/5, but none could say what moving them would do. Every control now explains its own effect.

THE INTERFACE

From blank canvas to *finished piece*.

Six surfaces — draw, generate, collect, share — designed as one continuous flow.



RESEARCH & INSIGHT

What stops us isn't imagination — *it's the skill to get it out.*

For people who love making pictures but can't yet make the one in their head — the idea always a few levels ahead of the hand.

- 01 · the gap

Imagination is free

Ideas come easily — a scene, a character, a feeling. Anyone can picture something they'd love to make.
- 02 · what's hard

Making is not

Turning the picture into a finished piece takes years of drawing skill — a barrier most people never cross.
- 03 · the shift

AI closes the gap

AI can carry the craft, so the limit becomes your imagination — not your hands.

"This may help those who can't draw/sketch well but know what they want their artpiece to look like."

— Cody · first reaction, before any explanation

the principle it leads to

Lower the barrier — but keep the person the **author**.

SketchMuse starts from the user's own sketch and uses AI only to finish what their hands couldn't — img2img, not text2img. The gate opens, and the person still holds the brush.

THE DIFFERENCE · same model, opposite user

Same img2img underneath — built for people, not engineers

Both start from your sketch. What happens next is where the tools split.

Other img2img tool · a raw model UI — you learn the controls by trial



SketchMuse · the same engine, wrapped in a real product



DESIGN DECISION MAKING

Four decisions behind the product

Every choice traded something away to keep the user in control.

- Start from a sketch, not text**

Input 1

img2img over text2img. The input is the user's own drawing.

source
Keep the person the author.

trade-off
You have to draw something first.
- Four results, not one**

Output 2

One sketch returns four variations — enough that one is usually worth keeping.

source
Keep the person the author.

trade-off
Six took too long. One left nothing to choose.
- Tune by experiment**

The model 3

Tested sampling method, CFG and denoising one by one for stable output.

source
Lower the barrier.

trade-off
Slow, manual testing, but each value earns its place.
- Jargon into plain words**

Controls 4

Renamed each control by what it does — CFG Scale became "Image Richness."


source
Lower the barrier.

trade-off
A clear name says what it is — not yet what it does.

THE LAYERS · simple by default, deep on demand

One system, revealed in three tiers

Complexity opens up only when someone asks for it. Nobody is pushed past the tier that fits.



THE MODEL · finding the default

Tuning by experiment, not guesswork

Over a thousand renders of the same apple, in batches of 2, 4, 6 and 12.

Sampling method · scheduler held constant

- apple · 01

LMS

detailed, photographic
- apple · 02

DPM++

flat, digital illustration
- apple · 03

LCM

vibrant but abstract
- apple · 04

DPM2

painterly, delicate
- apple · 05

DPM++2a

coloured-pencil, soft
- apple · 06

Heun

rich, true to the sketch

Schedule type · sampler fixed at Heun

- apple · 07

Exponential

flat, indistinct
- apple · 08

DDIM

ornamental, soft
- apple · 09

Simple

more detail, still thin
- apple · 10

Uniform

clean but plain
- apple · 11

Normal

glossy, speckled
- apple · 12

Beta

richest texture, colour

The verdict: Heun + Beta.

Across every test, this pairing gave the richest, most faithful result — so I locked it in as the default. Users get that quality with zero tuning.

HOW I BUILT IT · the full stack

From the model up to the product

One continuous build — from tuning the AI model to shipping the interface.

- A The model**

Tune the model

Tested samplers & schedulers → Heun + Beta, then studied CFG, denoising, prompts.
- B The product**

Design decisions

img2img over text, jargon into plain words, three tiers — the user stays the author.
- C The interface**




Design the UI / UX

Five pages, the flow between them, and the Standard / Advanced control system.
- D The code**

Build the front end

Hand-coded in HTML, CSS and JavaScript — canvas, interactions and every page, shipped solo.

From a fine-tuned diffusion model to a working interface — designed and built by one person.



TESTING · 5 users

What real users got stuck on

Five users, static screens, a questionnaire — no live session.




- Observed
1 The names read fine — that was the ceiling
I'd renamed the technical settings in plain words.
Rated for clarity, they scored 4.2/5 — three of five found them clear.
- Observed
2 The name only gave back the name
Asked what "Image Richness" changes, all five answered with the label itself.
None could say which way to move it, or that pushing it far enough makes the image worse.
- Observed
3 Four readings, none the same
Asked what "Sketch & Prompts Balance" does, four of five described something different — a before-and-after reveal, a sketch-to-text converter, a render button.
The one who read it correctly was the only engineer in the group.
- Observed
4 The prompt box was the hard part
Two of five named the prompt, not the sliders, as the step they hesitated on longest.

THE RESULT · what testing changed

What testing added

Nothing was renamed. Everything got explained.



STILL OPEN · what I'd test next

The fixes are in. The proof isn't.

Static screens found the misreadings. They never saw anyone stall.

01 · Understanding	02 · Behaviour	03 · Intent
Is Image Richness still misread now that it explains itself?	How long does a first prompt take with the style chips there?	Does anyone open the Advanced panel without being told?

The third tier only got a door this round — nobody has tried to find it yet.

One sketch in, a finished piece out — and a product *still learning* from the people who use it.

Turning doodles into finished art.

An AI platform that turns anyone's sketch into finished artwork.
No drawing skill required, and your original idea stays at the center.

DESIGNED & BUILT BY
Zoey Xiong
CONTACT
zoeyxiong0520@gmail.com